

```

#Question 1
library(gbm)

## Loaded gbm 2.1.8

library(plyr)
library(MLmetrics)

##
## Attaching package: 'MLmetrics'
## The following object is masked from 'package:base':
##
## Recall

library(Hmisc)

## Loading required package: lattice
## Loading required package: survival
## Loading required package: Formula
## Loading required package: ggplot2
##
## Attaching package: 'Hmisc'
## The following objects are masked from 'package:plyr':
##
## is.discrete, summarize
## The following objects are masked from 'package:base':
##
## format.pval, units

library(tidyverse);

## -- Attaching packages ----- tidyverse 1.3.0 --
## v tibble 3.0.3      v dplyr 1.0.0
## v tidyr 1.1.0      v stringr 1.4.0
## v readr 1.3.1      v forcats 0.5.0
## v purrr 0.3.4

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::arrange() masks plyr::arrange()
## x purrr::compact() masks plyr::compact()
## x dplyr::count() masks plyr::count()
## x dplyr::failwith() masks plyr::failwith()
## x dplyr::filter() masks stats::filter()
## x dplyr::id() masks plyr::id()
## x dplyr::lag() masks stats::lag()
## x dplyr::mutate() masks plyr::mutate()
## x dplyr::rename() masks plyr::rename()
## x dplyr::src() masks Hmisc::src()
## x dplyr::summarise() masks plyr::summarise()
## x dplyr::summarize() masks Hmisc::summarize(), plyr::summarize()

library(Rtsne);
data <- read_csv("datasets_26073_33239_weight-height.csv")

```

```
## Parsed with column specification:
## cols(
##   Gender = col_character(),
##   Height = col_double(),
##   Weight = col_double()
## )
```

```
data$Gender = factor(data$Gender,
                      levels = c("Male", "Female"),
                      labels = c(0, 1))
data$Gender <- as.character(data$Gender)
head(data$Gender)
```

```
## [1] "0" "0" "0" "0" "0" "0"
```

```
smp_size <- floor(0.75 * nrow(data))
```

```
## set the seed to make your partition reproducible
set.seed(100)
train_ind <- sample(seq_len(nrow(data)), size = smp_size)
```

```
train <- data[train_ind, ]
test <- data[-train_ind, ]
```

```
model <- gbm(Gender ~ Height + Weight, distribution="bernoulli", data=train,
              n.trees = 100,
              interaction.depth = 2,
              shrinkage = 0.1);
pred <- predict(model, newdata = test, type="response");
```

```
## Using 100 trees...
```

```
sum(pred>0.5)/nrow(test)
```

```
## [1] 0.506
```

1. The result of the last exercise was 0.464, the accuracy of the last exercise is close to this one which is 0.506.

```
data2 <- read_csv("datasets_38396_60978_characters_stats.csv")
```

```
## Parsed with column specification:
## cols(
##   Name = col_character(),
##   Alignment = col_character(),
##   Intelligence = col_double(),
##   Strength = col_double(),
##   Speed = col_double(),
##   Durability = col_double(),
##   Power = col_double(),
##   Combat = col_double(),
##   Total = col_double()
## )
```


```
data2
```

```
## # A tibble: 611 x 9
```


```
##   Name      Alignment Intelligence Strength Speed Durability Power Combat Total
##   <chr>      <chr>           <dbl>    <dbl> <dbl>      <dbl> <dbl>  <dbl> <dbl>
```

```
## 1 3-D Man    good      50      31      43      32      25      52      233
## 2 A-Bomb     good      38     100      17      80      17      64      316
## 3 Abe Sapi~  good      88      14      35      42      35      85      299
## 4 Abin Sur   good      50      90      53      64      84      65      406
## 5 Abominat~  bad       63      80      53      90      55      95      436
## 6 Abraxas    bad       88     100      83      99     100      56      526
## 7 Adam Mon~  good      63      10      12     100      71      64      320
## 8 Adam Str~  good       1       1       1       1       0       1       5
## 9 Agent 13   good       1       1       1       1       0       1       5
## 10 Agent Bob good      10       8      13       5       5      20      61
## # ... with 601 more rows
```


```
boxplot(data2[3])
```




```
boxplot(data2[4])
```




```
boxplot(data2[5])
```




```
boxplot(data2[6])
```




```
boxplot(data2[7])
```



```
boxplot(data2[8])
```



```
boxplot(data2[9])
```



1. Depending on the boxplot, there are not any irregularities and I don't need to filter out a subset of rows.

```
data2.pca <- prcomp(data2[,3:9])
data2.pca
```

```
## Standard deviations (1, ..., p=7):
## [1] 1.809591e+02 2.492497e+01 1.961815e+01 1.602658e+01 1.496131e+01
## [6] 1.451877e+01 2.352081e-14
##
## Rotation (n x k) = (7 x 7):
##           PC1      PC2      PC3      PC4      PC5
## Intelligence -0.1558138 -0.512410987 0.011954768 -0.10394268 0.62945701
## Strength     -0.1485702 0.574336433 0.387331365 -0.12549242 0.42880903
## Speed        -0.1155588 0.118684676 -0.190901802 0.88752932 -0.04624362
## Durability   -0.1834677 0.356326076 0.088248520 -0.27898376 -0.44316183
```

```
## Power      -0.1668084 -0.022455199 -0.764794152 -0.32552703 -0.09544381
## Combat     -0.1544233 -0.515763934  0.469773884 -0.01091879 -0.46054378
## Total      -0.9246422 -0.001282934  0.001612583  0.04266465  0.01287301
##              PC6      PC7
## Intelligence 0.40390193 -0.3779645
## Strength     -0.39440627 -0.3779645
## Speed        0.05841828 -0.3779645
## Durability   0.64381834 -0.3779645
## Power        -0.35892356 -0.3779645
## Combat       -0.36657410 -0.3779645
## Total        -0.01376538  0.3779645
```

```
summary(data2.pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation 180.9591 24.92497 19.61815 16.02658 14.9613 14.51877
## Proportion of Variance 0.9507 0.01804 0.01117 0.00746 0.0065 0.00612
## Cumulative Proportion 0.9507 0.96875 0.97992 0.98738 0.9939 1.00000
##              PC7
## Standard deviation 2.352e-14
## Proportion of Variance 0.000e+00
## Cumulative Proportion 1.000e+00
```


Q2:

2. According to the result of the principal component analysis, we need 1 component to get 85% of the variation in the data set.
3. We don't need to normalize these columns.
4. The total column is the total of the values in other numerical columns.
5. We shouldn't include that in the PCA, the largest principal components and the total column have the largest weighted values

```
library(ggfortify)
pca.plot <- autoplot(data2.pca, data = data2, colour = 'Alignment')
```

```
## Warning: 'select_()' is deprecated as of dplyr 0.7.0.
## Please use 'select()' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_warnings()' to see where this warning was generated.
```

```
pca.plot
```



6) The points of the 4 groups are clustered for the most part; however, the three points at the right of the graph may be outliers. The data does not appear to depart widely from multivariate normality.


```
library(Rtsne);
Labels <- data2$Alignment
data2$Alignment <- as.factor(data2$Alignment)
colors = rainbow(length(unique(data2$Alignment)))
names(colors) = unique(data2$Alignment)
tsne <- Rtsne(data2[,3:9], dims = 2, perplexity=30, verbose=TRUE, max_iter = 500, check_duplicates = FALSE)

## Performing PCA
## Read the 611 x 7 data matrix successfully!
## OpenMP is working. 1 threads.
## Using no_dims = 2, perplexity = 30.000000, and theta = 0.500000
## Computing input similarities...
## Building tree...
## Done in 0.07 seconds (sparsity = 0.193383)!
## Learning embedding...
## Iteration 50: error is 54.248133 (50 iterations in 0.10 seconds)
## Iteration 100: error is 49.383369 (50 iterations in 0.09 seconds)
## Iteration 150: error is 48.987868 (50 iterations in 0.08 seconds)
## Iteration 200: error is 48.816828 (50 iterations in 0.07 seconds)
## Iteration 250: error is 48.726268 (50 iterations in 0.08 seconds)
## Iteration 300: error is 0.713593 (50 iterations in 0.09 seconds)
## Iteration 350: error is 0.653791 (50 iterations in 0.09 seconds)
## Iteration 400: error is 0.635313 (50 iterations in 0.08 seconds)
## Iteration 450: error is 0.626903 (50 iterations in 0.07 seconds)
## Iteration 500: error is 0.620394 (50 iterations in 0.07 seconds)
```

```
## Fitting performed in 0.84 seconds.
exeTimeTsne<- system.time(Rtsne(data2[,3:9], dims = 2, perplexity=30, verbose=TRUE, max_iter = 500,check=TRUE))

## Performing PCA
## Read the 611 x 7 data matrix successfully!
## OpenMP is working. 1 threads.
## Using no_dims = 2, perplexity = 30.000000, and theta = 0.500000
## Computing input similarities...
## Building tree...
## Done in 0.06 seconds (sparsity = 0.193383)!
## Learning embedding...
## Iteration 50: error is 52.810262 (50 iterations in 0.09 seconds)
## Iteration 100: error is 49.291739 (50 iterations in 0.08 seconds)
## Iteration 150: error is 48.948248 (50 iterations in 0.08 seconds)
## Iteration 200: error is 48.793637 (50 iterations in 0.08 seconds)
## Iteration 250: error is 48.705934 (50 iterations in 0.09 seconds)
## Iteration 300: error is 0.724315 (50 iterations in 0.09 seconds)
## Iteration 350: error is 0.646306 (50 iterations in 0.08 seconds)
## Iteration 400: error is 0.629444 (50 iterations in 0.08 seconds)
## Iteration 450: error is 0.617899 (50 iterations in 0.07 seconds)
## Iteration 500: error is 0.610791 (50 iterations in 0.08 seconds)
## Fitting performed in 0.81 seconds.

plot(tsne$Y, t='n', main="tsne")
text(tsne$Y, labels=data2$Alignment, col=colors[data2$Alignment])
```



Q3:

The points within the individual clusters are highly similar to each other and close to points in other clusters. The same pattern likely holds in a high-dimensional original dataset. In the digits dataset, t-SNE didn't separate clusters of each digit class. In the context of alignment, these clusters represent the alignment with similar associated characters.

Q4: The plot produced by python is in the same directory that was named as 'tsne_2d'.

```
data3 = data2[,2:9]
data3
```

```
## # A tibble: 611 x 8
##   Alignment Intelligence Strength Speed Durability Power Combat Total
##   <fct>           <dbl>    <dbl> <dbl>    <dbl> <dbl>  <dbl> <dbl>
## 1 good             50      31    43      32    25    52    233
## 2 good             38     100    17      80    17    64    316
## 3 good             88      14    35      42    35    85    299
## 4 good             50      90    53      64    84    65    406
## 5 bad              63      80    53      90    55    95    436
## 6 bad              88     100    83      99   100    56    526
## 7 good             63      10    12     100    71    64    320
## 8 good              1        1    1        1     0     1     5
## 9 good              1        1    1        1     0     1     5
## 10 good            10        8    13         5     5    20     61
## # ... with 601 more rows
```

```
library(caret);
```


```
##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##   lift

## The following object is masked from 'package:survival':
##
##   cluster

## The following objects are masked from 'package:MLmetrics':
##
##   MAE, RMSE
```

```
control <- trainControl(method="repeatedcv", number=10, repeats=3)
set.seed(7)
modelGbm <- train(Alignment~., data=data3, method="gbm", na.action=na.exclude, trControl=control, verbose=0)
summary(modelGbm)
```



```
##           var  rel.inf
## Intelligence Intelligence 25.097294
## Total           Total 21.865455
## Speed           Speed 14.020442
## Strength        Strength 12.338157
## Combat          Combat 11.186908
## Power           Power 10.349115
## Durability      Durability 5.142628
```

```
modelGbm
```

```
## Stochastic Gradient Boosting
##
## 611 samples
## 7 predictor
## 3 classes: 'bad', 'good', 'neutral'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 545, 547, 547, 548, 546, 548, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  Accuracy  Kappa
##  1                   50      0.7132108  0.09383176
##  1                   100      0.7110793  0.10979519
##  1                   150      0.7044758  0.11139626
##  2                   50      0.7066983  0.11038054
##  2                   100      0.7001369  0.11865352
##  2                   150      0.6913041  0.10528900
##  3                   50      0.7089882  0.13808781
##  3                   100      0.6893035  0.10321287
##  3                   150      0.6931471  0.13589523
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
```

```
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 50, interaction.depth =
## 1, shrinkage = 0.1 and n.minobsinnode = 10.
```

```
set.seed(7)
data4 = data3
data4 = data4 %>% relocate(Alignment, .after = last_col())
data4 <- na.omit(data4)

control <- rfeControl(functions=rfFuncs, method="cv", number=10)
results <- rfe(data4[,1:7], data4[[8]], sizes=c(1:8),rfeControl=control)

print(results)
```

```
##
## Recursive feature selection
##
## Outer resampling method: Cross-Validated (10 fold)
##
## Resampling performance over subset size:
##
## Variables Accuracy   Kappa AccuracySD KappaSD Selected
##      1    0.6942 0.06296    0.03482 0.08919
##      2    0.7123 0.13521    0.03364 0.10764
##      3    0.7089 0.11555    0.04470 0.13908
##      4    0.7189 0.17224    0.04758 0.13336
##      5    0.7123 0.15976    0.03929 0.11315
##      6    0.7189 0.15603    0.04250 0.12807
##      7    0.7270 0.16442    0.02841 0.09814      *
##
## The top 5 variables (out of 7):
##      Intelligence, Combat, Durability, Power, Total
```

```
predictors(results)
```

```
## [1] "Intelligence" "Combat"      "Durability"   "Power"        "Total"
## [6] "Strength"     "Speed"
```

Q5: According to the results generated above, the best parameters are Intelligence, Combat, Durability, Power, Total.

Q6: k-fold Cross-validation is a resampling procedure used to evaluate machine learning models on a limited data sample.

The reasons why we need to use k-fold cross-validation are it is simple to understand and it is generally results in a less biased or less optimistic estimate of the model skill than other methods, such as a simple train/test split. We are able to 1)make predictions on all of our data using k-fold cross-validation. 2) get more metrics and draw important conclusion both about our algorithm and our data.3)work with dependent/grouped data. 4) Do Parameters Fine-Tuning

if we just report a single number for the accuracy of our model, this accuracy is more likely to be biased and inaccurate.

Q7: Recursive Feature Elimination, or RFE, is a popular feature selection algorithm. it is easy to configure and it is effective at selecting the best features.

RFE works by searching for a subset of features by starting with all features in the training dataset and

removing features until the desired number remains. In other words, RFE works as follows: fitting the given machine learning algorithm used in the core of the model, ranking features by importance, discarding the least important features, and re-fitting the model. This process is repeated until a specified number of features remains.